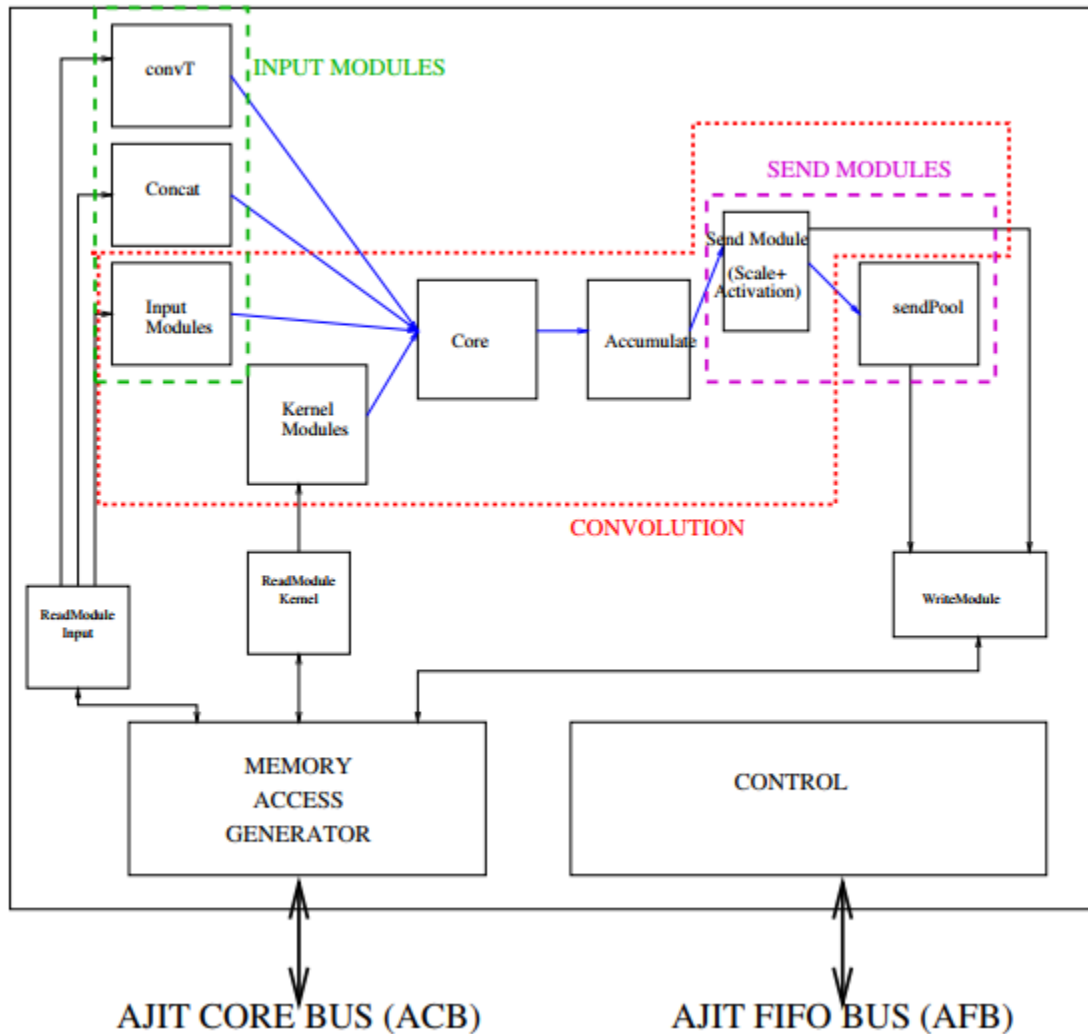


AI ML Accelerator Explanation

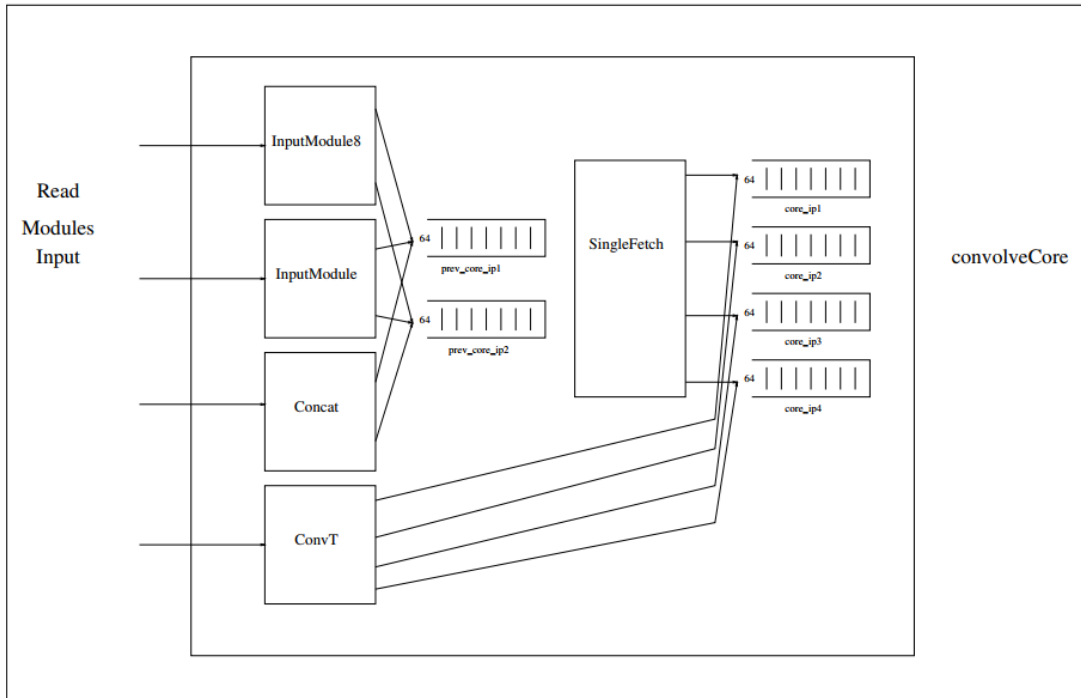
The overall block diagram of the AI ML accelerator designed by Aman Dhammani



Refer to the “An AI/ML Inference Engine Developed Using AHIR V2 Tools” thesis written by Aman Dhammani before going through the explanation.

This document explains chapter 3 from the thesis.

1. Fetching data from memory



- a. InputModule fetches 2 rows at a time. It has two submodules and the first submodule fetches odd numbered rows and stores it in `prev_core_ip1` and the second submodule fetches from even numbered rows and stores it in `prev_core_ip2`.

Ex: `prev_core_ip1` = [row 1, row 3, row 5]

`prev_core_ip2` = [row 2, row 4, row 6]

*Note: Here row "i" contains data of single row, single column and 8 channels(i.e indices = (i, j, chn:chn+8)), because each data width of row "i" is 64 bit and data type of single row, single column and single channel of input (i.e index=(i,j,chn)) is 8 bit

- b. SingleFetch

In first iteration, it writes odd row data to pipe 1 and even row data in pipe 2 From 2nd iteration onwards it writes odd row data to pipe 1 & 3 and even row data to pipe 2 and 4

Ex:

t = 0

`core_ip1` = [row 1]

`core_ip2` = [row 2]

```
core_ip3 = []  
core_ip4 = []
```

```
t = 1  
core_ip1 = [row 1, row 3]  
core_ip2 = [row 2, row 4]  
core_ip3 = [row 3]  
core_ip4 = [row 4]
```

```
t = 2  
core_ip1 = [row 1, row 3, row 5]  
core_ip2 = [row 2, row 4, row 6]  
core_ip3 = [row 3, row 5]  
core_ip4 = [row 4, row 6]
```

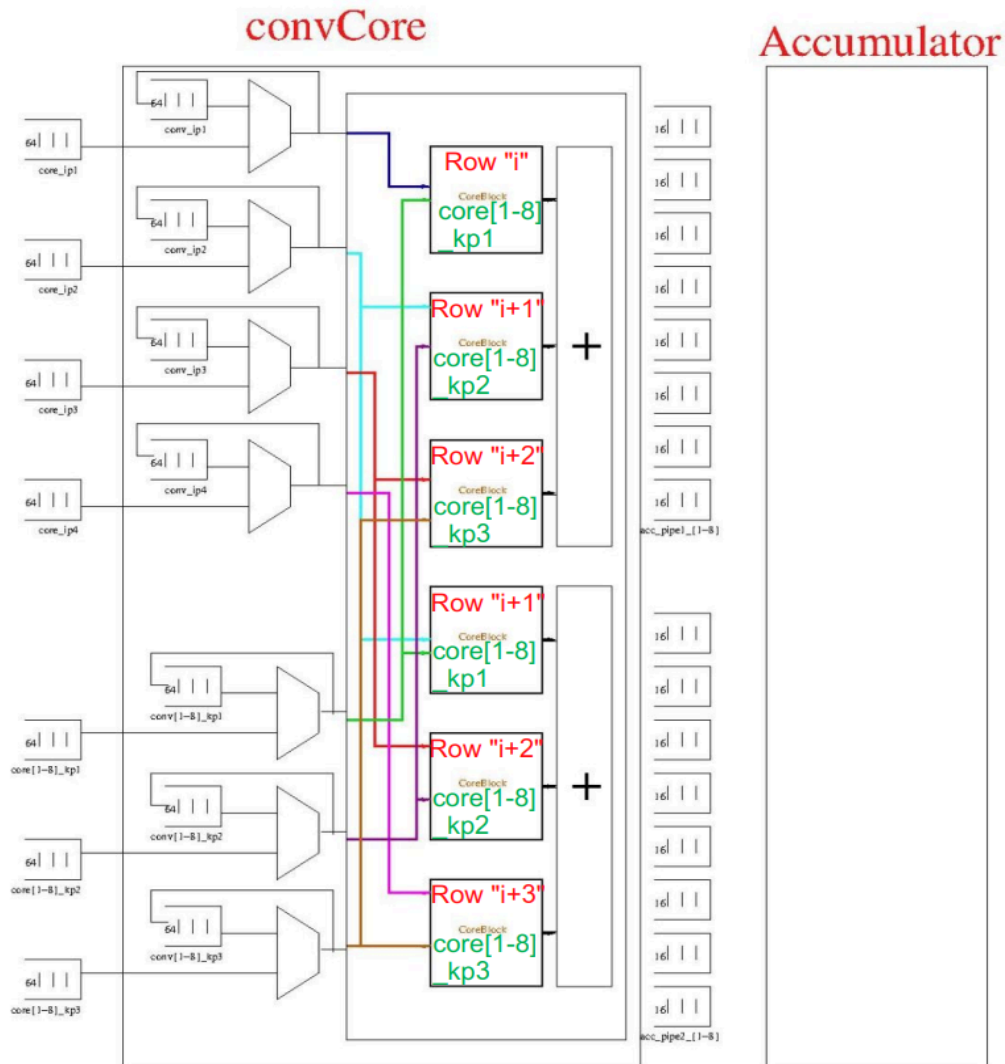
... and so on

c. Kernel Modules

- i. This set of modules is responsible for fetching the kernel data from the memory, and sending it to the core.
- ii. The kernel data is sent through three pipes, one for each row.
- iii. The kernel module makes calls to submodules to fetch kernels for different rows.
- iv. Since rk is at most three, the submodules need not fetch alternate rows.
- v. The fetch is done channel by channel, and is sent to pipes core<i> kp< j >, where < i > is the core number (corresponds to chn 1-8), and < j > is the row number, i.e. different for each submodule.

2. Convolution

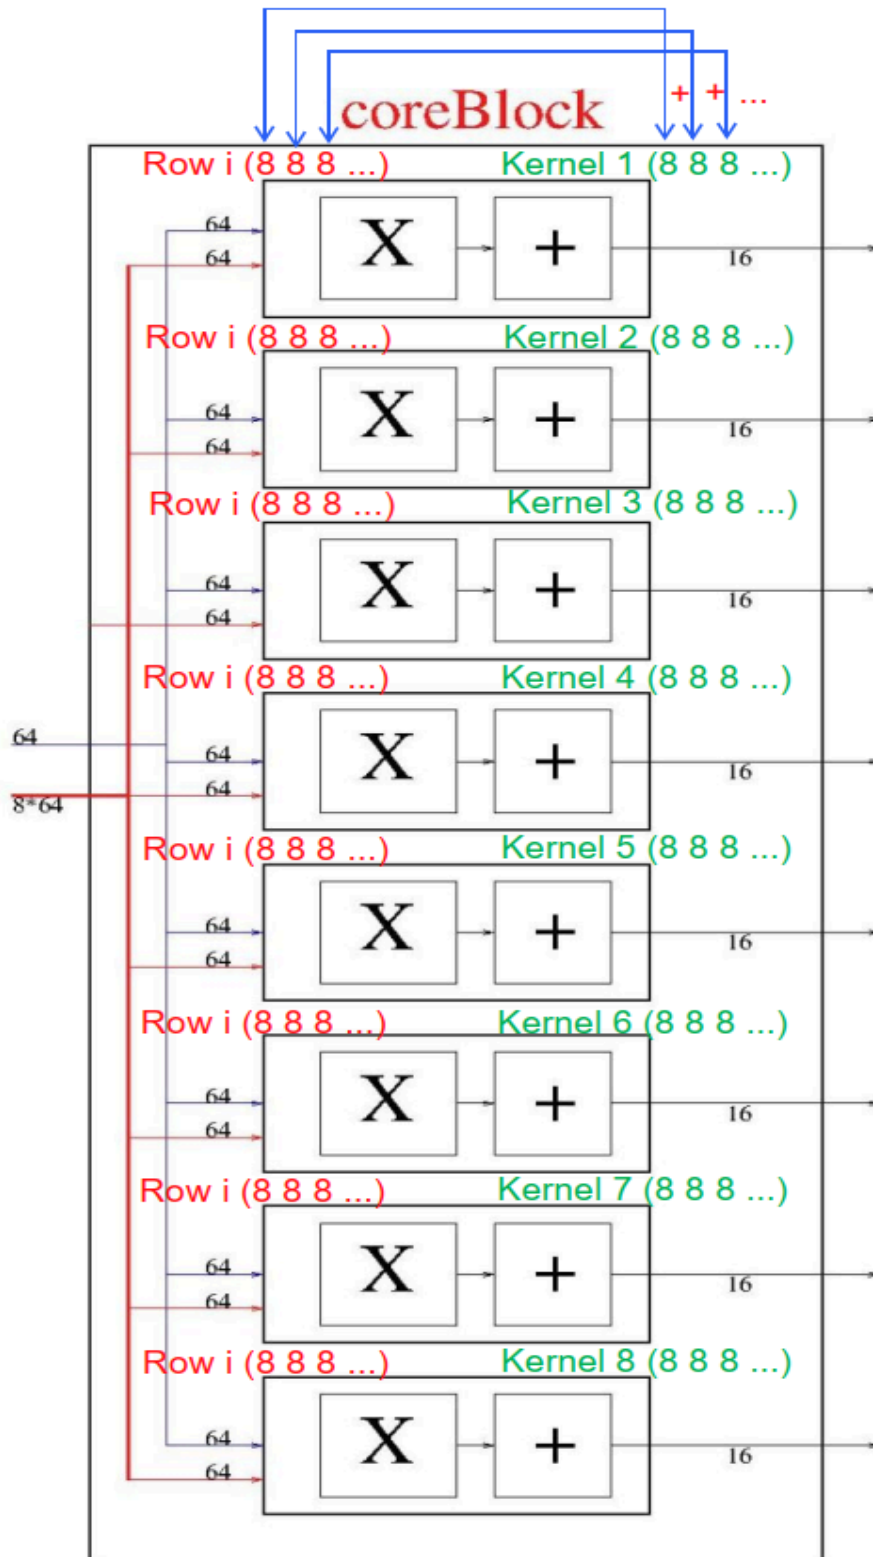
a. convolveCore



- There are 6 logical sections(i.e. CoreBlock) of 8*8 multipliers in the convolveCore.
- The first 3 logical sections calculate partial products for 1st output row and next 3 logical sections calculate partial products for 2nd output row.
- Each CoreBlock takes one input row of 64 bits and 8 kernels, each of 64 bits
- Each CoreBlock gives 8 partial products of 16 bits, one for each output channel, and the 8 partial products output of 3 CoreBlocks are added channel wise.

b. CoreBlock

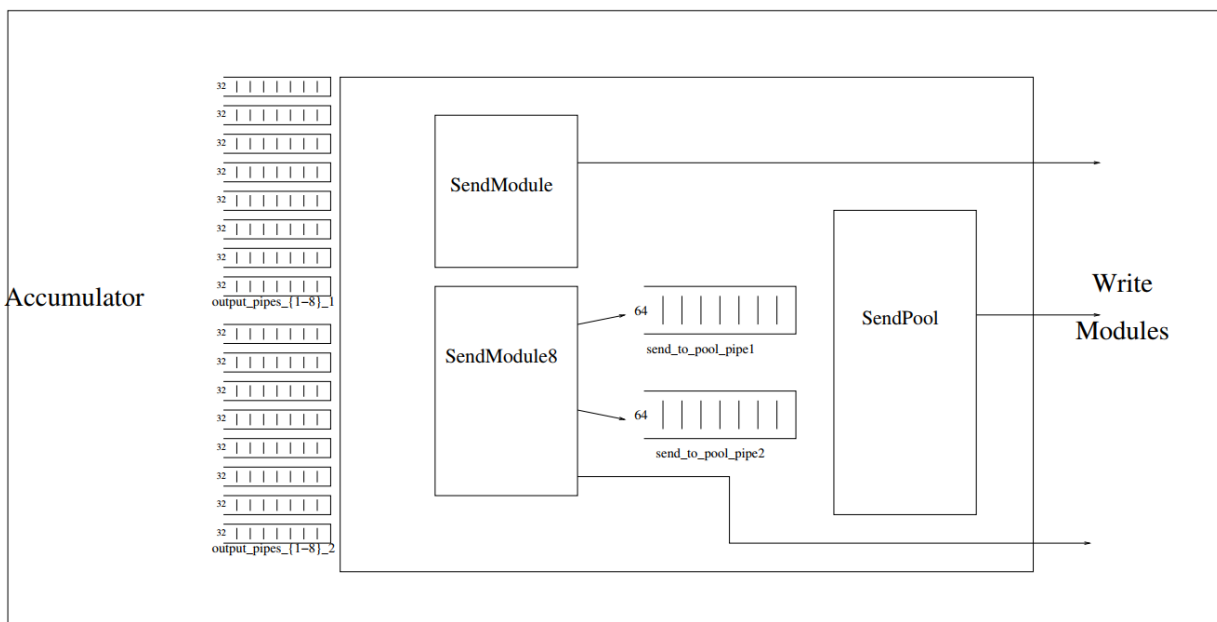
8 bits of input is multiplied
with 8 bits of kernel and
the product is added



c. Accumulator

- i. The accumulator reads the values and accumulates until count equals the parameter `acc_count`, which corresponds to the number of product of the number of channels and columns of the input required to get a output
- ii. After that, the counter and the accumulated value resets and waits for the next set of partial sums to produce the next output
- iii. The accumulator reads 16 output values across 2 rows and 8 output channels, and maintains partial sums that are 32 bit wide.
- iv. On completion of every accumulation, the data is forwarded to the send Module for further processing and write back to memory

3. Output Modules



- a. The output modules receives the data from the accumulator, and writes them back into the memory after applying the required operations.
- b. If the number of output channels is not a multiple of 8, it packs the data into worlds before writing it into memory.
- c. All the send modules are equipped to perform non linear activation on the output data which is received from the accumulator through two sets of 8 FIFOs each of 32 bit wide

4. Working

- a. The below part explains how the value of output in row “r”, column “c” and “chn_out” channels are calculated
- b. First 8 channels of row “r” and column “c” of input is convolved with first 8 channels of the row 1 and column 1 of the kernels 1-8
- c. Since convolution of 1 input with 1 kernel gives data for one output channel, hence using 8 kernels at a time gives values for 8 output channels.
- d. Next 8 input channels and 8 kernel channels are used until all the input and kernel channels are utilized
- e. Then similar convolution starts from first 8 channels of row “r+1” and column “c” of input and first 8 channels of row “2” and column “c” of kernels 1-8
- f. After iterating through all the rows, the same process is repeated for the next column of input and kernels.
- g. This convolution continues until it reaches the last column of the kernel.
- h. So above process from “c” - “h” is equivalent to convolution of input $[r:(r+r_k), c:(c+c_k), 1:\text{chn_in}]$ with 8 kernels $[1:r_k, 1:c_k, 1:\text{chn_in}]$ and produces the output values $[r, c, 1:8]$
- i. Now to calculate the output values in the channels [9:16], the above steps are repeated with same input but with kernels of 9 - 16
- j. This is continued until all the channels of row “r” and column “c” of the output is calculated
- k. In the actual implementation, row “r” and row “r+1” are simultaneously calculated.